

# Memory Printer and Adaptive Page Swapping in xv6

Umesh Kumar (2022CS11115)

Aditya Sahu (2022CS11113)

Operating Systems (COL331), Spring 2025

## Abstract

We present two major enhancements to the xv6 teaching kernel:

1. A *Memory Printer* facility that, upon user request, reports each active process's resident set size (RSS) in pages.
2. An *Adaptive Page Swapping* subsystem that automatically moves pages to and from disk under a fixed 4 MiB physical memory constraint, using a feedback-driven eviction policy parameterized by  $\alpha$  and  $\beta$ .

Detailed design, implementation snippets, and parameter analysis are provided.

## 1 Introduction

In modern operating systems, understanding per-process memory usage and gracefully handling memory overcommit are core responsibilities. We augment xv6 with:

- **Memory Printer:** On pressing `Ctrl+I`, the kernel prints “PID NUM\_PAGES” for each user process in RAM.
- **Adaptive Swapping:** When free pages fall below a threshold  $T_h$ , the kernel swaps out  $N_{pg}$  pages, then reduces  $T_h$  by factor  $(1 - \beta/100)$  and grows  $N_{pg}$  by  $(1 + \alpha/100)$ , all controlled via Makefile macros.

Our modifications touch the console driver, trap handler, page allocator, VM system, fs/mkfs, and a new `pageswap.c` module.

## 2 Memory Printer

### 2.1 Console Interrupt Hook

In `console.c`, we detect ASCII 9 (`Ctrl+I`) in the input loop:

```

1 // console.c
2 ...
3 case C('I'):
4     memory_printer();
5     break;
6 ...

```

## 2.2 Memory Printer Routine

Defined in `proc.c`, this function locks the process table, iterates active procs, and counts PTE\_P bits:

```

1 // proc.c
2
3 // Count resident pages for a process
4 static int
5 count_resident_pages(struct proc *p) {
6     int cnt = 0;
7     for (uint va = 0; va < p->sz; va += PGSIZE) {
8         pte_t *pte = walkpgdir(p->pgdir, (void*)va, 0);
9         if (pte && (*pte & PTE_P))
10             cnt++;
11     }
12     return cnt;
13 }
14
15 // Print header and per-PID counts
16 void
17 memory_printer(void) {
18     cprintf("Ctrl+I is detected by xv6\nPID NUM_PAGES\n");
19     acquire(&ptable.lock);
20     for (struct proc *p = ptable.proc; p < &ptable.proc[NPROC]; p++) {
21         if (p->pid >= 1 &&
22             (p->state==SLEEPING || p->state==RUNNABLE || p->state==
23              RUNNING)) {
24             cprintf("%d %d\n", p->pid, count_resident_pages(p));
25         }
26     }
27     release(&ptable.lock);
28 }

```

## 3 Adaptive Page Swapping

### 3.1 Disk Partitioning in mkfs.c

We reserve 800 slots of 8 sectors each immediately after the superblock:

```
1 // mkfs.c
2 // superblock layout adjustments:
3 sb.swapstart = xint(2);           // block 2
4 sb.swapsize  = xint(800 * 8);     // 800 slots * 8 blocks
5 sb.logstart  = xint(sb.swapstart + sb.swapsize);
```

We then write the updated superblock to sector 1.

### 3.2 Swap Slot Data Structure

In pageswap.h:

```
1 // pageswap.h
2 #define NUM_SLOTS 800
3 #define BLOCKS_PER_SLOT 8
4 #define PTE_S 0x200 // custom swapped-out flag
5
6 struct swapslot {
7     int page_perm; // original PTE flags
8     int is_free;   // 1 if slot is unused
9 };
```

### 3.3 Initialization at Boot

In fs.c, called from main():

```
1 // fs.c
2 void
3 init_swap(void) {
4     swapstart = 2; // as per superblock
5     for (int i = 0; i < NUM_SLOTS; i++) {
6         swap_slots[i].is_free = 1;
7         swap_slots[i].page_perm = 0;
8     }
9 }
```

### 3.4 Page Fault Handler

In trap.c, under case T\_PGFLT:

```
1 // trap.c
2 case T_PGFLT: {
```

```

3     uint va = rcr2();
4     struct proc *p = myproc();
5     pte_t *pte = walkpgdir(p->pgdir, (void*)va, 0);
6     if (pte && (*pte & PTE_S)) {
7         uint slot = *pte >> 12;
8         swap_in(p, va, slot);
9         return;
10    }
11    // otherwise: kill as usual
12    p->killed = 1;
13 } break;

```

### 3.5 Swapping Logic in pageswap.c

```

1 // pageswap.c
2
3 // Write 8 blocks of the page to disk
4 static void
5 write_to_swap(char *mem, int slot) {
6     for (int j = 0; j < BLOCKS_PER_SLOT; j++) {
7         struct buf *b = bread(ROOTDEV, swapstart + slot*8 + j);
8         memmove(b->data, mem + j*BSIZE, BSIZE);
9         bwrite(b);
10        brelse(b);
11    }
12 }
13
14 // Find and reserve a free slot
15 int
16 find_free_slot(void) {
17     for (int i = 0; i < NUM_SLOTS; i++)
18         if (swap_slots[i].is_free) {
19             swap_slots[i].is_free = 0;
20             return i;
21         }
22     panic("No free swap slots");
23 }
24
25 // Free a slot when page is either reloaded or process exits
26 void
27 free_swap_slot(uint slot) {
28     swap_slots[slot].is_free = 1;
29     swap_slots[slot].page_perm = 0;
30 }
31
32 // Swap out a single page

```

```

33 void
34 swap_out_one_page(void) {
35     // 1. Find victim process by max rss
36     struct proc *victim = 0; int maxrss = -1;
37     acquire(&ptable.lock);
38     for (struct proc *p = ptable.proc; p < &ptable.proc[NPROC]; p++) {
39         if ((p->state==RUNNING||p->state==RUNNABLE||p->state==SLEEPING)
40             && p->rss>maxrss) {
41             victim = p; maxrss = p->rss;
42         }
43     }
44     if (!victim) panic("swap_out: no victim");
45     // recompute rss
46     int cnt=0;
47     for (uint a = 0; a < victim->sz; a += PGSIZE)
48         if ((walkpgdir2(victim->pgdir,(void*)a,0)) &&
49             (*walkpgdir2(victim->pgdir,(void*)a,0)&PTE_P))
50             cnt++;
51     victim->rss = cnt;
52     release(&ptable.lock);
53
54     // 2. Select a page: first with PTE_P=1,PTE_A=0
55     uint va; pte_t *pte;
56     for (va=0; va<victim->sz; va+=PGSIZE) {
57         pte = walkpgdir2(victim->pgdir,(void*)va,0);
58         if (pte && (*pte & PTE_P) && !(*pte & PTE_A)) break;
59         if (pte) *pte &= ~PTE_A; // second chance
60     }
61     if (va>=victim->sz) panic("swap_out: no page");
62
63     // 3. Write out
64     int slot = find_free_slot();
65     char *mem = (char*)P2V(*pte & ~0xFFF);
66     swap_slots[slot].page_perm = *pte & 0xFFF;
67     write_to_swap(mem, slot);
68
69     // 4. Update PTE
70     *pte = (slot<<12) | PTE_S;
71     victim->rss--;
72     kfree(mem);
73 }
74
75 // Called by kalloc() when free pages <= Th
76 void
77 swap_out_if_needed(void) {
78     if (get_free_pages() <= Th) {
79         cprintf("Current Threshold = %d, Swapping %d pages\n", Th, Npg);

```

```

80     for (int i=0; i<Npg; i++) swap_out_one_page();
81     Th    = (Th*(100-BETA))/100;
82     Npg    = min(LIMIT, (Npg*(100+ALPHA))/100);
83 }
84 }
85
86 // Reload a swapped out page on fault
87 void
88 swap_in(struct proc *p, uint va, uint slot) {
89     swap_out_if_needed();
90     char *mem = kalloc();
91     read_from_swap(mem, slot);
92     pte_t *pte = walkpgdir2(p->pgdir, (void*)va, 0);
93     *pte = V2P(mem) | swap_slots[slot].page_perm | PTE_P;
94     p->rss++;
95     swap_slots[slot].is_free = 1;
96 }

```

## 4 Parameter Analysis

The adaptive controller is specified by:

$$\alpha \in [0, 100], \quad \beta \in [0, 100], \quad T_h(0) = 100, \quad N_{pg}(0) = 2, \quad \text{LIMIT} = 100.$$

Whenever the free-page count  $F \leq T_h$ , the kernel:

1. Logs “Current Threshold =  $T_h$ , Swapping  $N_{pg}$  pages.”
2. Calls `swap_out_one_page()` exactly  $N_{pg}$  times.
3. Updates

$$T_h \leftarrow \lfloor T_h (1 - \frac{\beta}{100}) \rfloor, \quad N_{pg} \leftarrow \min(\text{LIMIT}, \lfloor N_{pg} (1 + \frac{\alpha}{100}) \rfloor).$$

### 4.1 Effect of $\alpha$ (Growth Factor)

- **Low  $\alpha$  (e.g. 0–10):**

- Batch size  $N_{pg}$  grows slowly (or not at all).
- *Pros*: Minimizes I/O burst sizes, preserves working set longer.
- *Cons*: Under-eviction on sustained high pressure  $\rightarrow$  repeated page faults.

- **Moderate  $\alpha$  (20–40):**

- Balanced growth: each eviction round frees more pages than the last, adapting to severity.

- Typically yields good throughput in mixed workloads.
- **High  $\alpha$  (50–100):**
  - Aggressive growth (doubling or more).
  - *Pros*: Quickly drains memory under pathological pressure.
  - *Cons*: Large I/O spikes, potential to evict pages that will soon be reused (“cache thrash”).

## 4.2 Effect of $\beta$ (Decay Factor)

- **Low  $\beta$  (0–10):**
  - Threshold  $T_h$  remains near its initial value.
  - Evictions occur rarely but in large batches (once triggered).
  - Risks *longer stalls* before any swapping kicks in.
- **Moderate  $\beta$  (10–30):**
  - Gradual lowering of  $T_h$ , leading to more frequent but smaller swap bursts.
  - Smooths out memory pressure handling.
- **High  $\beta$  (40–100):**
  - Rapid decay of  $T_h$ —subsequent rounds trigger almost immediately, potentially over-evicting.
  - May cause repeated small I/O writes, lowering effective throughput.

## 4.3 Tuning Trade-Offs

This controller behaves like a proportional feedback loop:

$$\Delta N_{pg} \propto \alpha, \quad \Delta T_h \propto -\beta.$$

One tunes  $(\alpha, \beta)$  to balance:

- *Fault Latency*: Time a process waits on a page-in (smaller  $N_{pg}$  and lower  $\beta$  help).
- *I/O Efficiency*: Amortization of each disk write (larger  $N_{pg}$  and higher  $\beta$  help).
- *Working-Set Preservation*: Risk of evicting soon-to-be-used pages (lower  $\alpha$ ,  $\beta$ ).

Empirical tuning in xv6 with  $\alpha = 25, \beta = 10$  provides smooth behavior under typical academic workloads.